



Case File CE-001

Suspicious token behavior without the hype

An anonymized forensic case file showing how an AI-assisted engine separates market-abuse evidence from unsafe trading claims.

ANONYMIZED

RESEARCH-ONLY

NO TRADE SIGNAL

NO PUBLIC ACCUSATION

82**Behavior score**

Strong anomaly

66.0%**Circularity**

Wash-like

Yes**Sell observed**

Not honeypot wall

No**Trade-ready**

Policy blocked

Forensic thesis

The engine found a coordinated, circular pattern. Then it refused to call it alpha because the outcome economics failed after friction. That refusal is the product.



Executive Signal, Not Hype

The short version

- The case was built from local raw Swap, Sync and Transfer evidence.
- Behavioral anomaly score reached 82/100 from dense timing, repeatable swaps, wallet concentration and transfer/liquidity context.
- Behavioral scoring does not require official protocol docs. Identity proof is separate and remains required before mapping, labels or client-facing outputs.
- Outcome testing rejected the long thesis: no positive shadow edge after conservative friction assumptions.

Public-safe conclusion

This is not a public accusation against any specific project or wallet. It is a demonstration of a forensic workflow: detect abnormal coordination, test whether it is tradable, and block unsafe conclusions.



Evidence Stack

Evidence	Count	Meaning
Raw swaps	50	Exact market interactions
Sync events	50	Reserve/liquidity context
Transfers	37	Token flow context
Unique txs	32	Repeatability
Wallets	23	Participant breadth

Why this is useful

- A case file is not a single indicator. It is a chain of corroborating observations.
- The engine keeps raw behavior separate from identity attribution and trading decisions.
- This makes the output suitable for research desks, listing-risk teams and suspicious-volume screening.



Coordination Timeline

Observed swap span: 2.9167 minutes. Max swaps in 5m/20 blocks: 50.



Reading: the AI could detect the abnormal pattern, but even earlier thresholds did not create an exploitable long edge after friction.



Circularity & Wash-Like Flow

66.0%

Circularity

high_circularity

16


Buy flows

Target buys

34

Sell flows


Target sells



- High circularity is consistent with wash-like volume or repeated internal flow.
- Hold/freeze proxy was 44.2572%, not enough for clean accumulation.
- This is valuable risk intelligence, but not a long-entry trigger.



Participant & Funder Graph



- Funder graph status: needs_funder_graph_collection.
- Fresh-like wallets observed: 3; missing wallet state rows: 20.
- Public conclusion: graph hints exist, but attribution remains blocked until stronger evidence exists.



Backtest Economics

0.25%

Best raw MFE

Theoretical upside

-9.75%

After friction

Reality check

28

Post-T0 sells

Sellability observed

The decisive point

- This was not rejected because the token was impossible to sell.
- It was rejected because suspicious behavior did not become profitable after realistic frictions.
- That makes the engine useful as a risk filter before it is useful as an execution system.



Policy Gate

NO SIGNAL

NO TRADE

NO MAPPING

NO LABEL

NO OPT-IN

NO ACCUSATION

Blocked surfaces

- No client signal because the economic replay is negative after friction.
- No mapping or label because identity/source proof remains separate from behavioral scoring.
- No trade because policy requires repeatability, controls, liquidity, risk limits and explicit opt-in.
- No public accusation because this is an anonymized proof-of-capability report.



Business Use Case

Who would care?

- Funds: avoid entering tokens where volume is mostly circular or artificial.
- Listing desks: flag suspicious token-market behavior before exposure.
- Forensic teams: prioritize clusters and funder paths without manually combing every pool.
- Market-data platforms: offer a manipulation-risk API, not another generic price alert.

Positioning

The strongest business model may not be automated trading first. It may be market-abuse intelligence: an API that says when a token's apparent activity is suspicious, circular, late, or economically unsafe.



Conclusion

The first credible demo is a refusal.

CE-001 shows an engine that detects suspicious behavior and then refuses the unsafe conclusion. That is more credible than a bot claiming every anomaly is alpha.

Recommended public framing

Publish the anonymized thread. Offer the PDF privately to serious researchers/funds. Avoid naming a token, wallet or protocol until source-backed attribution exists.

- This report demonstrates forensic capability.
- It does not disclose proprietary heuristics fully.
- It invites demand discovery: if people ask for an API, the market is speaking.