



Case File CE-001

Suspicious token behavior without the hype

An anonymized forensic case file showing how an AI-assisted engine separates market-abuse evidence from unsafe trading claims.

ANONYMIZED

RESEARCH-ONLY

NO TRADE SIGNAL

NO PUBLIC ACCUSATION

82

Behavior score
Strong anomaly

66.0%

Circularity
Wash-like

Yes

Sell observed
Not honeypot wall

No

Trade-ready
Policy blocked

Forensic thesis

The engine found a coordinated, circular pattern. Then it refused to call it alpha because the outcome economics failed after friction. That refusal is the product.

The short version

- The case was built from local raw Swap, Sync and Transfer evidence — no third-party claims, no social signals, no price feed involved in detection.
- Behavioral anomaly score reached **82/100** from dense timing, repeatable swaps, wallet concentration and transfer/liquidity context.
- Behavioral scoring does not require official protocol docs. Identity proof is separate and remains required before mapping, labels or client-facing outputs.
- Outcome testing rejected the long thesis: **no positive shadow edge after conservative friction assumptions.**

Public-safe conclusion

This is not a public accusation against any specific project or wallet. It is a demonstration of a forensic workflow: detect abnormal coordination, test whether it is tradable, and block unsafe conclusions.

What the engine actually saw

EVIDENCE	COUNT	MEANING
Raw swaps	50	Exact market interactions
Sync events	50	Reserve / liquidity context
Transfers	37	Token flow context
Unique txs	32	Repeatability
Wallets	23	Participant breadth

Why a chain, not an indicator

- A case file is not a single indicator. It is a chain of corroborating observations.
- The engine keeps raw behavior separate from identity attribution and trading decisions.
- This makes the output suitable for research desks, listing-risk teams and suspicious-volume screening.

Three minutes of dense swaps

Observed swap span: 2.9167 minutes. Max swaps in 5m/20 blocks: 50.




Each S-value is a detection-sensitivity threshold. Lower thresholds detect the pattern earlier; S82 is the anomaly score finally recorded. The "lead" is how many minutes before the price peak an entry at that threshold would have given. Missing values at S70/S82 are honest: at those sensitivities the entry window no longer resolves.

Reading


The engine could detect the abnormal pattern, but even earlier thresholds did not create an exploitable long edge after friction.

Two thirds of the volume circles back

**66.0%****Circularity**
high_circularity**16****Buy flows**
Target buys**34****Sell flows**
Target sells

- High circularity is consistent with wash-like volume or repeated internal flow.
- Hold/freeze proxy was 44.26% — not enough for clean accumulation.
- This is valuable risk intelligence, but not a long-entry trigger.

Hints exist, attribution blocked



- Funder graph status: **needs_funder_graph_collection** — the graph layer is a hint, not an attribution.
- Fresh-like wallets observed: 3; missing wallet state rows: 20.
- Public conclusion: graph hints exist, but attribution remains blocked until stronger evidence exists.

Where the thesis died

0.25%

Best raw MFE
Theoretical upside

-9.75%

After friction
Reality check

28

Post-T0 sells
Sellability observed



The friction-adjusted outcome is ~39× worse than the theoretical upside. A friction-blind scanner would have called this token "actionable".

The decisive point

This was not rejected because the token was impossible to sell — 28 post-T0 sell rows were observed. It was rejected because suspicious behavior did not become profitable after realistic frictions. That makes the engine useful as a risk filter before it is useful as an execution system.

What the engine refuses to do

NO SIGNAL

NO TRADE

NO MAPPING

NO LABEL

NO OPT-IN

NO ACCUSATION

Blocked surfaces

- **No client signal** — because the economic replay is negative after friction.
- **No mapping or label** — because identity/source proof remains separate from behavioral scoring.
- **No trade** — because policy requires repeatability, controls, liquidity, risk limits and explicit opt-in.
- **No public accusation** — because this is an anonymized proof-of-capability report.

Why this page exists

A forensic engine is defined by what it refuses as much as by what it detects. The gate list is the product's liability boundary — and the reason it can be shown publicly.

Who would care, and why

Who would care?

Funds — avoid entering tokens where volume is mostly circular or artificial.

Listing desks — flag suspicious token-market behavior before exposure.

Forensic teams — prioritize clusters and funder paths without manually combing every pool.

Market-data platforms — offer a manipulation-risk API, not another generic price alert.

Positioning

The strongest business model may not be automated trading first. It may be **market-abuse intelligence**: an API that says when a token's apparent activity is suspicious, circular, late, or economically unsafe.

The first credible demo is a refusal

CE-001 shows an engine that detects suspicious behavior and then refuses the unsafe conclusion. That is more credible than a bot claiming every anomaly is alpha.

Recommended public framing

Publish the anonymized thread. Offer the PDF privately to serious researchers and funds. Avoid naming a token, wallet or protocol until source-backed attribution exists.

- This report demonstrates forensic capability.
- It does not disclose proprietary heuristics fully.
- It invites demand discovery: if people ask for an API, the market is speaking.



ANONYMOUS MARKET INTELLIGENCE — CE-001 — END OF FILE